

NYC Bike Share Analysis

Ed Powell

March 3, 2019

```
#install.packages("dplyr")
#install.packages("ggplot2")
#install.packages("lubridate")
#install.packages("geosphere")
#install.packages("xts")
#install.packages("randomForest")
#install.packages("caret")
#install.packages("lattice")
#install.packages("e1071")
#install.packages("splitstackshape")
#install.packages("tree")
#install.packages("devtools")
#install.packages("knitr")
#install.packages("markdown")
#install.packages("rmarkdown")

library(dplyr)
library(ggplot2)
library(lubridate)
library(geosphere)
library(xts)
library(randomForest)
library(caret)
library(e1071)
library(splitstackshape)
library(tree)
library(devtools)
library(knitr)
library(markdown)
library(rmarkdown)
nycbikedata <- read.csv("C://Users/ep927/Documents/NewYorkCity/CleanedBikeData.csv", header = TRUE)

nycbikedata$Start_Year <- strptime(nycbikedata$Start_Time, "%m/%d/%Y %H:%M")

nycbikedata$Year <- year(nycbikedata$Start_Year)
nycbikedata$Month <- month(nycbikedata$Start_Year)
nycbikedata$DayOfWeek <- wday(nycbikedata$Start_Year, label = TRUE, abbr = TRUE)
```

```
nycbikedata$Qtr <- quarter(nycbikedata$Start_Year)
nycbikedata$DayOfWeek <- ifelse(wday(nycbikedata$Start_Year)==1,7,wday(nycbikedata$Start_Year)-1)
nycbikedata$Age <- (nycbikedata$Year - nycbikedata$Birth_Year)
nycbikedata$distance <- distHaversine(nycbikedata[,6:7], nycbikedata[,10:11],r=6378137) / 1609.344
nycbikedata$RideInMin <- round(nycbikedata$Trip_Duration / 60 ,digits = 0)
```

```
str(nycbikedata)
```

```
## 'data.frame': 802870 obs. of 23 variables:
## $ Trip_Duration : int 61 61 61 61 61 61 61 61 61 61 ...
## $ Start_Time : Factor w/ 498008 levels "1/1/2016 0:02",...: 20
729 199918 247126 308008 341766 441798 416524 425976 427339 496816 ...
## $ Stop_Time : Factor w/ 498759 levels "1/1/2016 0:08",...: 20
871 200195 247410 308231 342159 442373 417035 426491 427839 497548 ...
## $ Start_Station_ID : int 3185 3185 3185 3185 3185 3267 3270 3272 3
186 3186 ...
## $ Start_Station_Name : Factor w/ 61 levels "5 Corners Library",...: 11
11 11 11 11 45 33 32 24 24 ...
## $ Start_Station_Longitude: num -74 -74 -74 -74 -74 ...
## $ Start_Station_Latitude: num 40.7 40.7 40.7 40.7 40.7 ...
## $ End_Station_ID : int 3186 3186 3186 3186 3186 3214 3272 3270 3
185 3185 ...
## $ End_Station_Name : Factor w/ 183 levels "11 Ave & W 41 St",...: 76
76 76 76 76 62 86 87 38 38 ...
## $ End_Station_Longitude : num -74 -74 -74 -74 -74 ...
## $ End_Station_Latitude : num 40.7 40.7 40.7 40.7 40.7 ...
## $ Bike_ID : int 24424 24491 24720 24675 24454 26176 24686
24666 26162 26253 ...
## $ User_Type : Factor w/ 2 levels "Customer","Subscriber": 2
2 2 2 2 2 2 2 2 ...
## $ Birth_Year : int 1983 1981 1983 1983 1981 1968 1990 1984 1
960 1992 ...
## $ Gender : int 1 1 1 1 1 1 1 1 1 1 ...
## $ Start_Year : POSIXlt, format: "2016-01-05 20:17:00" "2016-0
3-30 09:08:00" ...
## $ Year : num 2016 2016 2016 2016 2016 ...
## $ Month : num 1 3 5 6 7 8 8 8 8 9 ...
## $ DayOfWeek : num 2 3 3 1 1 1 2 7 1 5 ...
## $ Qtr : int 1 1 2 2 3 3 3 3 3 3 ...
## $ Age : num 33 35 33 33 35 48 26 32 56 24 ...
## $ distance : num 0.134 0.134 0.134 0.134 0.134 ...
## $ RideInMin : num 1 1 1 1 1 1 1 1 1 1 ...
```

```
summary(nycbikedata)
```

```
## Trip_Duration Start_Time Stop_Time
## Min. : 61.0 11/2/2016 8:17 : 14 10/4/2018 8:28 : 13
```

```

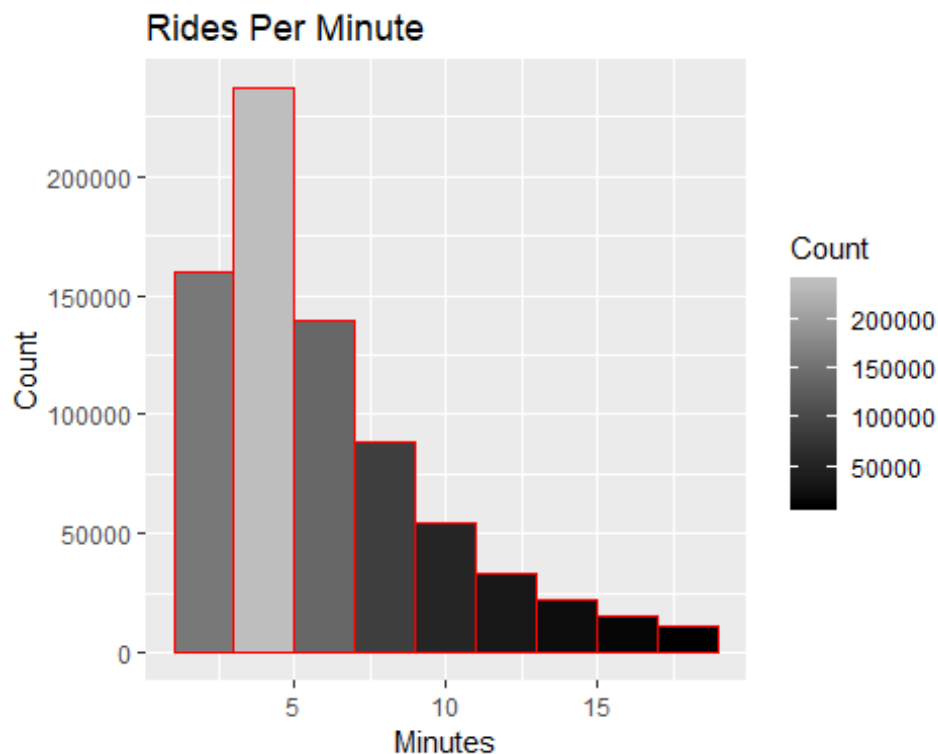
## 1st Qu.: 229.0    10/4/2018 8:22 :    13    10/10/2018 8:57:    12
## Median : 333.0    10/10/2018 8:41:    12    10/9/2018 8:46 :    12
## Mean   : 476.7    10/12/2018 8:28:    12    6/29/2018 8:25 :    12
## 3rd Qu.: 534.0    6/8/2018 8:02 :    12    6/5/2018 8:45 :    12
## Max.   :40968.0    8/10/2018 8:28 :    12    9/28/2016 8:37 :    12
##                                     (Other) :802795    (Other) :802797
## Start_Station_ID      Start_Station_Name Start_Station_Longitude
## Min.   :3183      Grove St PATH : 99088    Min.   : -74.10
## 1st Qu.:3187      Hamilton Park : 50761    1st Qu.: -74.05
## Median :3203      Exchange Place: 50260    Median : -74.04
## Mean   :3235      Sip Ave       : 48254    Mean   : -74.05
## 3rd Qu.:3267      Newport PATH  : 36234    3rd Qu.: -74.04
## Max.   :3694      Newark Ave   : 25244    Max.   : -74.03
##                                     (Other) :493029
## Start_.Station_Latitude End_Station_ID      End_Station_Name
## Min.   :40.69      Min.   : 79    Grove St PATH :127115
## 1st Qu.:40.72      1st Qu.:3186    Exchange Place: 62630
## Median :40.72      Median :3202    Hamilton Park : 47848
## Mean   :40.72      Mean   :3230    Sip Ave       : 43900
## 3rd Qu.:40.73      3rd Qu.:3220    Newport PATH  : 38593
## Max.   :40.75      Max.   :3694    Newark Ave   : 23980
##                                     (Other) :458804
## End_Station_Longitude End_Station_Latitude Bike_ID
## Min.   : -74.10      Min.   :40.68      Min.   :14529
## 1st Qu.: -74.05      1st Qu.:40.72      1st Qu.:26157
## Median : -74.04      Median :40.72      Median :26314
## Mean   : -74.05      Mean   :40.72      Mean   :27852
## 3rd Qu.: -74.04      3rd Qu.:40.73      3rd Qu.:29587
## Max.   : -73.95      Max.   :40.81      Max.   :35009
##
## User_Type      Birth_Year      Gender
## Customer : 9959    Min.   :1940    Min.   :1.000
## Subscriber:792911 1st Qu.:1975    1st Qu.:1.000
##                                     Median :1983    Median :1.000
##                                     Mean   :1980    Mean   :1.224
##                                     3rd Qu.:1988    3rd Qu.:1.000
##                                     Max.   :2002    Max.   :2.000
##
## Start_Year      Year      Month
## Min.   :2016-01-01 00:02:00    Min.   :2016    Min.   : 1.000
## 1st Qu.:2016-11-23 11:37:00    1st Qu.:2016    1st Qu.: 5.000
## Median :2017-09-19 12:18:30    Median :2017    Median : 7.000
## Mean   :2017-09-08 10:29:15    Mean   :2017    Mean   : 7.209
## 3rd Qu.:2018-06-22 20:08:30    3rd Qu.:2018    3rd Qu.:10.000
## Max.   :2018-12-31 23:25:00    Max.   :2018    Max.   :12.000
##
## DayOfWeek      Qtr      Age      distance
## Min.   :1.000    Min.   :1.000    Min.   :16.00    Min.   :0.08678
## 1st Qu.:2.000    1st Qu.:2.000    1st Qu.:30.00    1st Qu.:0.38596
## Median :4.000    Median :3.000    Median :35.00    Median :0.55040

```

```
## Mean :3.685 Mean :2.737 Mean :36.86 Mean :0.65634
## 3rd Qu.:5.000 3rd Qu.:4.000 3rd Qu.:42.00 3rd Qu.:0.84295
## Max. :7.000 Max. :4.000 Max. :78.00 Max. :7.54271
##
## RideInMin
## Min. : 1.000
## 1st Qu.: 4.000
## Median : 6.000
## Mean : 7.945
## 3rd Qu.: 9.000
## Max. :683.000
##
```

Graphs for some of the datat

```
ggplot(data=nycbikedata, aes(nycbikedata$RideInMin)) +
  labs(title = "Rides Per Minute") +
  labs(x="Minutes", y="Count") +
  geom_histogram(breaks=seq(1,20, by =2),
                 col="red",
                 aes(fill=..count..)) +
  scale_fill_gradient("Count", low = "black", high = "gray")
```




```
ggplot(data=nycbikedata, aes(nycbikedata$Age)) +
  labs(title = "Riders Age") +
  labs(x="Age", y="Count") +
  geom_histogram(breaks=seq(15,80, by =2),
```

```

col="red",
aes(fill=..count..) +
scale_fill_gradient("Count", low = "orange", high = "purple"
)


```




```

ggplot(data=nycbikedata, aes(nycbikedata$Month)) +
  labs(title = "Riders By Month") +
  labs(x="Month", y="Count") +
  geom_histogram(breaks=seq(1,12,by = 2),
    col="red",
    aes(fill=..count..) +
    scale_fill_gradient("Count", low = "orange", high = "purple"
  )

```




```
ggplot(nycbikedata, aes(Age, RideInMin)) +  
  labs(title = "Line is the mean Ride") +  
  geom_point(color = "grey") +  
  stat_summary(fun.y = "mean", geom = "line", size = 1, linetype = "dashed")
```



```
ggplot(nycbikedata, aes(Age, distance)) +  
  geom_point(color = "green") +  
  labs(title = "The dashed line represents the mean distance ridden by Age") +  
  stat_summary(fun.y = "mean", geom = "line", size = 1, linetype = "dashed")
```

The dashed line represents the mean distance ridden by



#Selecting the fields to be used in predicting Gender using RandomForset. Since Gender is a int I need to convert it to a Factor

```
df <- select(nycbikedata, Age, Gender, Year, Month, DayOfWeek, distance, RideInMin, Start_Station_ID, End_Station_ID)
```

```
small.sample <- stratified(df, "Gender", size=10000)
```

```
small.sample$Gender <- as.factor(small.sample$Gender)
```

```
table(small.sample$Gender)
```

```
##
```

```
##      1      2
```

```
## 10000 10000
```

```
str(small.sample)
```

```
## Classes 'data.table' and 'data.frame': 20000 obs. of 9 variables:
```

```
## $ Age : num 46 34 32 29 31 42 29 37 47 32 ...
```

```
## $ Gender : Factor w/ 2 levels "1","2": 1 1 1 1 1 1 1 1 1 1 ...
```

```
## $ Year : num 2017 2016 2018 2018 2017 ...
```

```
## $ Month : num 4 9 11 7 2 8 11 1 7 10 ...
```

```
## $ DayOfWeek : num 3 2 2 2 1 7 3 4 3 2 ...
```

```
## $ distance : num 0.424 0.386 1.01 1.898 0.583 ...
```

```
## $ RideInMin : num 4 2 8 11 6 44 4 7 6 5 ...
```

```
## $ Start_Station_ID: int 3195 3279 3214 3194 3187 3211 3194 3183 3194 327
```



```

0 ...
## $ End_Station_ID : int 3194 3186 3202 3183 3214 3213 3195 3276 3195 318
5 ...
## - attr(*, ".internal.selfref")=<externalptr>

#Data Partition
#ind = Independent Variable. The data will be split 70/30
#rf = Random Forest

set.seed(123)
ind <- sample(2, nrow(small.sample), replace=TRUE, prob = c(0.7, 0.3))
train <- small.sample[ind==1,]
test <- small.sample[ind==2,]

#Random Forset - Prior to Tuning the model

set.seed(111)
rf <- randomForest(Gender~., data = train)
print(rf)

##
## Call:
## randomForest(formula = Gender ~ ., data = train)
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 2
##
## OOB estimate of error rate: 31.8%
## Confusion matrix:
##      1      2 class.error
## 1 4730 2318  0.3288876
## 2 2159 4873  0.3070250

attributes(rf)

## $names
## [1] "call"           "type"           "predicted"
## [4] "err.rate"       "confusion"      "votes"
## [7] "oob.times"      "classes"        "importance"
## [10] "importanceSD"   "localImportance" "proximity"
## [13] "ntree"          "mtry"           "forest"
## [16] "y"              "test"           "inbag"
## [19] "terms"
##
## $class
## [1] "randomForest.formula" "randomForest"

#Prediction and Confusion Matrix - Train Data
#Pred1 = Predication 1
pred1 <- predict(rf, train)
confusionMatrix(pred1, train$Gender)

```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    1    2
```

```
##           1 7041   36
```

```
##           2    7 6996
```

```
##
```

```
##           Accuracy : 0.9969
```

```
##           95% CI : (0.9959, 0.9978)
```

```
##           No Information Rate : 0.5006
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.9939
```

```
##           McNemar's Test P-Value : 1.955e-05
```

```
##
```

```
##           Sensitivity : 0.9990
```

```
##           Specificity : 0.9949
```

```
##           Pos Pred Value : 0.9949
```

```
##           Neg Pred Value : 0.9990
```

```
##           Prevalence : 0.5006
```

```
##           Detection Rate : 0.5001
```

```
##           Detection Prevalence : 0.5026
```

```
##           Balanced Accuracy : 0.9969
```

```
##
```

```
##           'Positive' Class : 1
```

```
##
```

```
#Prediction and Confusion with Matrix - Test Data
```

```
#Pred2 = Prediction 2
```

```
pred2 <- predict(rf, test)
```

```
confusionMatrix(pred2, test$Gender)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    1    2
```

```
##           1 1989  932
```

```
##           2  963 2036
```

```
##
```

```
##           Accuracy : 0.6799
```

```
##           95% CI : (0.6678, 0.6918)
```

```
##           No Information Rate : 0.5014
```

```
##           P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.3598
```

```
##           McNemar's Test P-Value : 0.4907
```

```
##
```


```
##           Sensitivity : 0.6738
```

```
##           Specificity : 0.6860
```

```
##          Pos Pred Value : 0.6809
##          Neg Pred Value : 0.6789
##          Prevalence : 0.4986
##          Detection Rate : 0.3360
##          Detection Prevalence : 0.4934
##          Balanced Accuracy : 0.6799
##
##          'Positive' Class : 1
##
```

#Error Rate

```
plot(rf)
```



#Tune my try


#train <- as.data.frame(train) will remove Error in randomForest.default(x, y, mtry = mtryStart, ntree = ntreeTry, : length of response must be the same as predictors

```
train <- as.data.frame(train)
t <- tuneRF(train[,-2], train[,2],
  stepFactor = 0.5,
  plot=TRUE,
  improve = 0.05)
```

```
## mtry = 2 OOB error = 34.62%
```

```
## Searching left ...
```

```
## mtry = 4      OOB error = 33.9%
## 0.0207222 0.05
## Searching right ...
## mtry = 1      OOB error = 39.14%
## -0.1306935 0.05
```



```
#Tune my model
# using the information from mtry I used that information to tune my model.

rf <- randomForest(Gender~., data = train,
  ntree = 400,
  mtry = 4,
  importance = TRUE,
  proximity = TRUE)

print(rf)


##
## Call:
## randomForest(formula = Gender ~ ., data = train, ntree = 400,      mtry =
4, importance = TRUE, proximity = TRUE)
##           Type of random forest: classification
##           Number of trees: 400
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 31.78%
```

```
## Confusion matrix:
##      1      2 class.error
## 1 4659 2389    0.3389614
## 2 2085 4947    0.2965017

rf$confusion

##      1      2 class.error
## 1 4659 2389    0.3389614
## 2 2085 4947    0.2965017

hist(treesize(rf),
     main = "Number Of Nodes for the Tree",
     col = "green")
```



```
#After modeled tuned

pred1 <- predict(rf, train)
confusionMatrix(pred1,train$Gender)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      1      2
##              1 7048    12
##              2     0 7020
##
##              Accuracy : 0.9991
```

```
##              95% CI : (0.9985, 0.9996)
##      No Information Rate : 0.5006
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9983
##      McNemar's Test P-Value : 0.001496
##
##      Sensitivity : 1.0000
##      Specificity : 0.9983
##      Pos Pred Value : 0.9983
##      Neg Pred Value : 1.0000
##      Prevalence : 0.5006
##      Detection Rate : 0.5006
##      Detection Prevalence : 0.5014
##      Balanced Accuracy : 0.9991
##
##      'Positive' Class : 1
##
```

#After Modeled tuned


```
pred2 <- predict(rf, test)
confusionMatrix(pred2, test$Gender)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    1    2
##           1 1981  848
##           2  971 2120
##
##              Accuracy : 0.6927
##              95% CI : (0.6808, 0.7045)
##      No Information Rate : 0.5014
##      P-Value [Acc > NIR] : < 2e-16
##
##              Kappa : 0.3854
##      McNemar's Test P-Value : 0.00423
##
##      Sensitivity : 0.6711
##      Specificity : 0.7143
##      Pos Pred Value : 0.7002
##      Neg Pred Value : 0.6859
##      Prevalence : 0.4986
##      Detection Rate : 0.3346
##      Detection Prevalence : 0.4779
##      Balanced Accuracy : 0.6927
##
##      'Positive' Class : 1
##
```

```
#Importance of Variables
```

```
varImpPlot(rf,  
           sort = TRUE,  
           n.var = 6,  
           main = "Top Six - Variable Importance")
```

Top Six - Variable Importance




```
varUsed(rf)
```

```
## [1] 231706 84733 196827 159184 179307 128370 138755 143979
```

```
partialPlot(rf, train, RideInMin, "1")
```

Partial Dependence on RideInMin



```
partialPlot(rf,train,RideInMin,"2")
```

Partial Dependence on RideInMin

